

File system implementation

Tran, Van Hoai

Faculty of Computer Science & Engineering
HCMC University of Technology

E-mail: hoai@hcmut.edu.vn
(*partly based on slides of Le Thanh Van*)

Outline

- 1 File system structure
- 2 File system implementation
- 3 Directory implementation
- 4 Allocation methods
- 5 Free-space management
- 6 Efficiency and performance
- 7 Recovery
- 8 Network file system

Outline

- 1 File system structure
- 2 File system implementation
- 3 Directory implemention
- 4 Allocation methods
- 5 Free-space management
- 6 Efficiency and performance
- 7 Recovery
- 8 Network file system

Design principles

Two design problems


- 1** Defining how the file system should look like to the user
files, attributes, file operations, directory structure,...
- 2** Creating algorithms/data structures to map logical file system to physical secondary-storage devices

Design principles

Two design problems

- 1 Defining how the file system should look like to the user
files, attributes, file operations, directory structure,...
- 2 Creating algorithms/data structures to map logical file system to physical secondary-storage devices

- **Basic file system:** to read/write physical blocks on disk
- **File-organization module:** files, logical/physical blocks
- **Logical file system:** all metadata information (no actual data)



Outline

- 1 File system structure
- 2 File system implementation**
- 3 Directory implemention
- 4 Allocation methods
- 5 Free-space management
- 6 Efficiency and performance
- 7 Recovery
- 8 Network file system

On-disk and in-memory structures

On-disk

- Boot control block (per volume): first block of volume, contain boot information of an OS
- Volume control block (per volume): partition details (#blocks, block size, free-blocks, free-FCBs)
- Directory structure: file organization
- per-file FCB

In-memory


- Mount table
- Directory structure cache
- System-wide open-file table
- Per-process open-file table
- Buffer holding file-system blocks

File Control Block


file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

File open/read

File open



File read



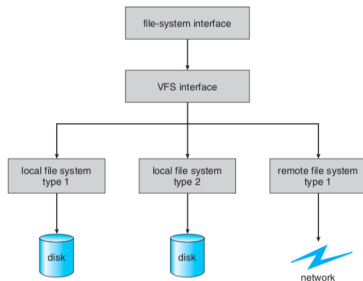
Virtual file system (VFS)

How to support multiple file systems in a directory structure ?

Virtual file system (VFS)

How to support multiple file systems in a directory structure ?

- Virtual File Systems (VFS) provide an **object-oriented** way of implementing file systems
- VFS allows the same API to be used for different types of file systems



API

API (open, close, read, write, mmap) defined on objects (inode object, file object, superblock object, dentry object)

Outline

- 1 File system structure
- 2 File system implementation
- 3 Directory implementation**
- 4 Allocation methods
- 5 Free-space management
- 6 Efficiency and performance
- 7 Recovery
- 8 Network file system

Directory implementation

- **Linear list:** a **linear** list of file names with pointers to data blocks
 - **Simple** to program
 - **Time-consuming** to execute
- **Hash table:** a linear list with hash data structure
 - **Small** directory search time
 - Collisions - situations where 2 file names hash to same location
 - **Fixed** size

Outline

- 1 File system structure
- 2 File system implementation
- 3 Directory implementation
- 4 Allocation methods**
- 5 Free-space management
- 6 Efficiency and performance
- 7 Recovery
- 8 Network file system


Types of allocation

Direct-access nature of disks gives us flexibility in the implementation of files.

- An allocation method refers to how disk blocks are allocated for files
- Types:
 - Contiguous allocation
 - Linked allocation
 - Indexed allocation

Contiguous allocation (1)

- Each file occupies a set of **contiguous blocks** on disk
- **Simple** – only starting location (block #) and length (number of blocks) are required.
- **no** (or **small**) disk head movement for random or sequential access
- **External** fragmentation (as in dynamic storage-allocation problem)
- Files **cannot grow**




Contiguous allocation (2)

Extent-based systems

- Many newer file systems (I.e. Veritas File System) use a modified contiguous allocation scheme
- Extent-based file systems allocate disk blocks in **extents**
- An extent is a **contiguous block** of disks. Extents are allocated for file allocation. A file consists of one or more extents


Linked allocation

- Each file is a **linked list** of disk blocks: blocks may be scattered anywhere on the disk.
- **Simple** - need only starting address
- Free-space management system → **no** waste of space (no external fragmentation)
- **No random** access
- **Low** reliability - especially when losing a block in a chain (no pointer)
- FAT - File Allocation Table (by MS-DOS, OS/2)




Linked allocation

FAT



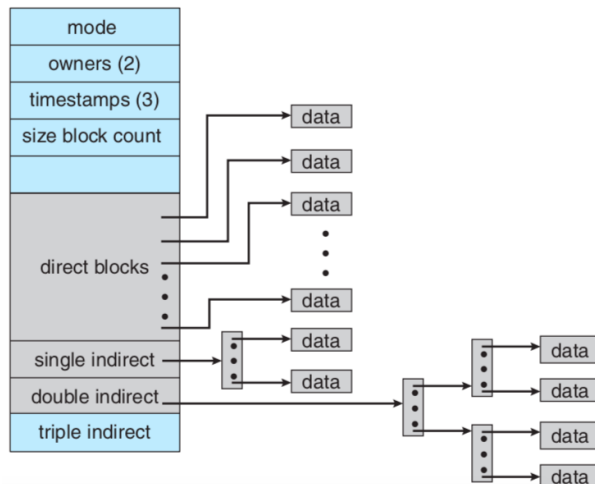
Indexed allocation

- Bring all pointers together into **index block**
- **random** access
- Dynamic access **without external** fragmentation
- **Overhead** of index block.
How index block adapt with different file sizes ?
 - Linked scheme
 - Multilevel index
 - Combined scheme



Indexed allocation

UNIX inode



Outline

- 1 File system structure
- 2 File system implementation
- 3 Directory implementation
- 4 Allocation methods
- 5 Free-space management**
- 6 Efficiency and performance
- 7 Recovery
- 8 Network file system

Free-space management

Free-space management is done by **free-space list** which keeps track disk space

- **Bit vector** (or bit map):



$$\text{bit}_i = \begin{cases} 0 & \text{block } i \text{ occupied} \\ 1 & \text{block } i \text{ free} \end{cases}$$

- **Block number calculation**

(Number of bits per word) \times
(Number of 0-valued words) +
offset of first 1 bit

- Copy of bit vector in memory and in disk must be consistent
- Keeping whole bit-vector in memory is expensive


Block size = 2^{12} bytes

Disk size = 2^{30} bytes (or 1GB)

$\Rightarrow n = 2^{30}/2^{12} = 2^{18}$ bits (or 32K bytes)

Linked list

- A linked list to link together all the free disk blocks




Linked list

- A linked list to link together all the free disk blocks

Other free-space management

- Grouping
- Counting
- Space maps





Outline

- 1 File system structure
- 2 File system implementation
- 3 Directory implementation
- 4 Allocation methods
- 5 Free-space management
- 6 Efficiency and performance**
- 7 Recovery
- 8 Network file system

Improving performance by cache

- Disk controllers has its own on-board cache
- OS maintains a separate section of main memory for **buffer cache** to store blocks **used again shortly**
- **Page cache**: using virtual memory techniques to cache file data (**in pages**, **not in file-system oriented blocks**)
⇒ Unified virtual memory (in Solaris, linux, Windows)



Outline

- 1 File system structure
- 2 File system implementation
- 3 Directory implementation
- 4 Allocation methods
- 5 Free-space management
- 6 Efficiency and performance
- 7 Recovery**
- 8 Network file system

Improving performance by cache

- **Consistency checking**: compares data in directory structure with data blocks on disk, and tries to fix inconsistencies
 - Use system programs to back up data from disk to another storage device (floppy disk, magnetic tape)
 - Recover lost file or disk by restoring data from backup
- **Log structured** (or journaling) file systems record each update to the file system as a transaction
 - All transactions are written to a log. A transaction is considered committed once it is written to the log. However, the file system may not yet be updated
 - The transactions in the log are asynchronously written to the file system. When the file system is modified, the transaction is removed from the log
 - If the file system crashes, all remaining transactions in the log must still be performed

Outline

- 1 File system structure
- 2 File system implementation
- 3 Directory implemention
- 4 Allocation methods
- 5 Free-space management
- 6 Efficiency and performance
- 7 Recovery
- 8 Network file system**

NFS architecture

